An Efficient Minibatch Acceptance Test for Metropolis-Hastings

Anonymous Author 1
Unknown Institution 1

Anonymous Author 2
Unknown Institution 2

Anonymous Author 3
Unknown Institution 3

Abstract

Markov chain Monte Carlo (MCMC) methods have many applications in machine learning. We are particularly interested in their application to modeling very large datasets, where it is impractical to perform Metropolis-Hastings tests on the full data. Previous work on reducing the cost of Metropolis-Hastings tests yield variable data consumed per sample, with only constant factor reductions versus using the full dataset for each sample. Here we present a method that can be tuned to provide arbitrarily small batch sizes, by adjusting either proposal step size or temperature. Our approach uses the natural noise present in minibatch likelihood estimates to furnish the randomness in a Metropolis-Hastings test. Our test uses the noise-tolerant Barker acceptance test with a novel additive correction variable. The resulting test can be combined with minibatch proposals to yield updates with the same complexity as a simple SGD update. In this paper we derive the test, analyze its performance, discuss its implementation, and present several experiments. We demonstrate several order-of-magnitude speedup over previous work, and show for the first time that small expected minibatch sizes are possible.

1 INTRODUCTION

Markov chain Monte Carlo (MCMC) sampling is a powerful method for computation on intractable distributions. We are interested in large dataset applications, where the goal is to sample a posterior distribution $p(\theta \mid x_1, \ldots, x_N)$ of parameter θ for large N. The Metropolis-Hastings method (M-H) generates sample

candidates from a proposal distribution q which is in general different from the target distribution p, and decides whether to accept or reject based on an acceptance test. The acceptance test is usually a Metropolis test [Metropolis et al., 1953, Hastings, 1970].

Many state-of-the-art machine learning methods, and deep learning in particular, are based on minibatch updates (such as SGD) to a model. Minibatch updates produce many improvements to the model for each pass over the dataset, and have high sample efficiency. In contrast, conventional M-H requires calculations over the full dataset to produce a new sample. Recent results from Korattikara et al. [2014] and Bardenet et al. [2014] perform approximate (bounded error) acceptance tests using subsets of the full dataset. The amount of data consumed for each test varies significantly from one minibatch to the next. By contrast, Maclaurin and Adams [2014], Bardenet et al. [2015] perform exact tests but require a lower bound on parameter distribution across its domain. The amount of data reduction depends on the accuracy of this bound, and such bounds are only available for relatively simple distributions.

Here we derive a new test which incorporates the variability in minibatch statistics as *a natural part of the test* and requires less data per iteration than prior work. We use a Barker test function [Barker, 1965], which makes our test naturally error tolerant. The idea of using a noise-tolerant Barker's test function was suggested but not explored empirically in Bardenet et al. [2015] section 6.3. But the asymptotic test statistic CDF and the Barker function are different, which leads to fixed errors for the approach in Bardenet et al. [2015]. Here, we show that the difference between the distributions can be corrected with an additive random variable. This leads to a test which is fast, and whose error can be made arbitrarily small.

Our test is applicable when the variance (over data samples) of the log acceptance probability is small enough (less than 1). It's not clear at first why this quantity should be bounded, but we will show that it is "natural" for well-specified models running Metropolis-Hastings sampling with optimal proposals [Roberts and Rosenthal, 2001] on a full dataset. When we reduce the

amount of data for the test, the variance goes up. We have to reduce variance in one of several ways. Either:

- Increase the temperature of the target distribution. Log likelihoods scale as $1/T$, and so the variance of the likelihood ratio will vary as $1/T^2$. Our model is no longer well-specified (we are doing inference at a temperature different from that assumed during data generation), but higher temperature can be advantageous for parameter exploration.
- Increase the minibatch size when needed. Log acceptance variance scales as $1/k$ vs the minibatch size k. Our test is adaptive like earlier works, but unlike them, the distribution of minibatch size is Gaussian, not long-tailed. Increased minibatch size also reduces the error rate for the test.
- For continuous distributions, reduce the proposal step size and variance compared to an optimal proposal. The variance of the log acceptance probability scales as the square of proposal step size.

It is worth discussing at this point the typical goals of M-H sampling on large datasets. By the Bernstein-von Mises Theorem, the posterior distribution for a Bayesian inference task has variance that scales inversely with N. Simply sampling from it is one application, but an efficient proposal [Roberts and Rosenthal, 2001] has similar variance to the target and will diffuse to it extremely slowly. For applications to neural networks or models where the posterior is multimodal [Choromanska et al., 2015], samplers will likely get trapped in one of the modes. A common solution is to anneal the sampler, running first at high temperatures to flatten the likelihood landscape. This in turn reduces the variance of the log acceptance probability and allows our test to be applied. Our samples can cover the search space densely with small steps rather than taking a few sparse steps towards an optimum. In this mode, Metropolis-Hastings can be used in similar fashion to Stochastic Gradient Descent. The goal in SGD is to make gradual progress to a posterior mode with each step, taking small steps so that the cumulative displacement has progressively lower variance.

The contributions of this paper are as follows:

- We develop a new, more efficient (in samples per test) minibatch acceptance test with quantifiable error bounds. The test uses a novel additive correction variable to implement a Barker test based on minibatch mean and variance.
- We compare performance of our new test and prior approaches on several datasets. We demonstrate orders of magnitude improvements in efficiency (measured as data consumed per test), and that it does not suffer from long-tailed minibatch sizes.

2 PRELIMINARIES

In the Metropolis-Hastings method [Gilks and Spiegelhalter, 1996, Brooks et al., 2011], a difficult-to-compute probability distribution $p(\theta)$ is sampled using a Markov chain $\theta_1, \ldots, \theta_n$. The sample θ_{t+1} at time $t+1$ is generated using a candidate θ' from a (simpler) proposal distribution $q(\theta' \mid \theta_t)$, filtered by an acceptance test. The acceptance test is usually a Metropolis test. The Metropolis test has acceptance probability:

$$\alpha(\theta_t, \theta') = \frac{p(\theta')q(\theta_t \mid \theta')}{p(\theta_t)q(\theta' \mid \theta_t)} \wedge 1 \qquad (1)$$

where $a \wedge b$ denotes $\min(a, b)$. With probability $\alpha(\theta_t, \theta')$, we accept θ' and set $\theta_{t+1} = \theta'$, otherwise set $\theta_{t+1} = \theta_t$. The test is often implemented with an auxiliary random variable $u \sim \mathcal{U}(0, 1)$ with a comparison $u < \alpha(\theta_t, \theta')$; here, $\mathcal{U}(a, b)$ denotes the uniform distribution on the interval $[a, b]$. For simplicity, we drop the subscript t for the current sample θ_t and denote it as θ.

The acceptance test guarantees detailed balance, which means $p(\theta)p(\theta' \mid \theta) = p(\theta')p(\theta \mid \theta')$, where $p(\theta' \mid \theta)$ is the probability of a transition from state θ to θ'. Here, $p(\theta' \mid \theta) = q(\theta' \mid \theta)\alpha(\theta, \theta')$. This condition, together with ergodicity, guarantees that the Markov chain has a unique stationary distribution $\pi(\theta) = p(\theta)$. For Bayesian inference, the target distribution is $p(\theta \mid x_1, \ldots, x_N)$. The acceptance probability is now:

$$\alpha(\theta, \theta') = \frac{p_0(\theta') \prod_{i=1}^{N} p(x_i \mid \theta')q(\theta \mid \theta')}{p_0(\theta) \prod_{i=1}^{N} p(x_i \mid \theta)q(\theta' \mid \theta)} \wedge 1 \qquad (2)$$

where $p_0(\theta)$ is the prior. Computing samples this way requires all N data points, but this is very expensive for large datasets. To address this challenge, Korattikara et al. [2014], Bardenet et al. [2014] perform approximate Metropolis-Hasting tests using sequential hypothesis testing. Each iteration, they start with a small minibatch and test whether θ' should be accepted based on approximating $u < \alpha(\theta, \theta')$. If the approximate test cannot decide with sufficient confidence, the minibatch size is increased and the test repeats. This process continues until a decision. The bounds depend on either an asymptotic Central Limit Theorem [Korattikara et al., 2014] or a concentration bound [Bardenet et al., 2014]. The latter requires direct bounds on the log likelihood ratio, which for general distributions requires knowing $p(x_i \mid \theta)$ and $p(x_i \mid \theta')$ for all N samples. In addition, while both methods show useful reductions in the number of samples required, they suffer the drawback of resolving small log likelihood ratio differences between the minibatch and full batch versions. We discuss a worst-case scenario in Section 2.1.

Following Bardenet et al. [2014], we write the test $u <$

$\alpha(\theta, \theta')$ equivalently as $\Lambda(\theta, \theta') > \psi(u, \theta, \theta')$, where[1]

$$\Lambda(\theta, \theta') = \sum_{i=1}^{N} \log\left(\frac{p(x_i|\theta')}{p(x_i|\theta)}\right)$$

$$\text{and} \quad \psi(u, \theta, \theta') = \log\left(u\frac{q(\theta'|\theta)p_0(\theta)}{q(\theta|\theta')p_0(\theta')}\right). \quad (3)$$

To reduce computational effort, an unbiased estimate of $\Lambda(\theta, \theta')$ based on a minibatch can be used:

$$\Lambda^*(\theta, \theta') = \frac{N}{b}\sum_{i=1}^{b}\log\left(\frac{p(x_i|\theta')}{p(x_i|\theta)}\right) \quad (4)$$

Finally, it will be convenient for our analysis to define $\Lambda_i(\theta, \theta') = N\log(\frac{p(x_i|\theta')}{p(x_i|\theta)})$. Thus, $\Lambda(\theta, \theta')$ is the mean of $\Lambda_i(\theta, \theta')$ over the entire dataset, and $\Lambda^*(\theta, \theta')$ is the mean of $\Lambda_i(\theta, \theta')$ over its minibatch.

Since minibatches contains randomly selected samples x_i, the values Λ_i are i.i.d. random variables[2]. By the Central Limit Theorem, we expect $\Lambda^*(\theta, \theta')$ to be approximately Gaussian. The acceptance test then becomes a statistical test of the hypothesis that $\Lambda(\theta, \theta') > \psi(u, \theta, \theta')$ by establishing that $\Lambda^*(\theta, \theta')$ is substantially larger than $\psi(u, \theta, \theta')$.

2.1 A Worst-Case Gaussian Example

Let $x_1, \ldots, x_N$ be i.i.d. $\mathcal{N}(\theta, 1)$ with known variance $\sigma^2 = 1$ and (unknown) mean $\theta = 0.5$. We use a uniform prior on θ. The log likelihood ratio is

$$\Lambda^*(\theta, \theta') = N(\theta' - \theta)\left(\frac{1}{b}\sum_{i=1}^{b}x_i - \theta - \frac{\theta' - \theta}{2}\right) \quad (5)$$

which is normally distributed over selection of the Normal samples x_i. Since the x_i have unit variance, their mean has variance $1/b$, and the variance of $\Lambda^*(\theta, \theta')$ is $\sigma^2(\Lambda^*) = (\theta' - \theta)^2 N^2/b$. In order to pass a hypothesis test that $\Lambda > \psi$, there needs to be a large enough gap (several $\sigma(\Lambda^*)$) between $\Lambda^*(\theta, \theta')$ and $\psi(u, \theta, \theta')$.

The posterior is a Gaussian centered on the sample mean μ, and with variance $1/N$ (i.e., $\mathcal{N}(\mu, 1/N)$). In one dimension, an efficient proposal distribution has the same variance as the target distribution [Roberts and Rosenthal, 2001], so we use a proposal based on $\mathcal{N}(\theta, 1/N)$. It is symmetric $q(\theta' \mid \theta) = q(\theta \mid \theta')$, and since we assumed a uniform prior, $\psi(u, \theta, \theta') = \log u$. Our worst-case scenario is specified in Lemma 1.

---

[1]Our definitions differ from those in Bardenet et al. [2014] by a factor of N to simplify our analysis later.

[2]The analysis assumes sampling with replacement although implementations on typical large datasets will approximate this by sampling without replacement.

Lemma 1. *For the model in Section 2.1, there exists a fixed (independent of N) constant c such that with probability $\geq c$ over the joint distribution of (θ, θ', u), the tests from Korattikara et al. [2014], Bardenet et al. [2014] consume all N samples.*

Proof. See Supplementary Material, Section A. □

Similar results can be shown for other distributions and proposals by identifying regions in product space (θ, θ', u) such that the hypothesis test needs to separate nearly-equal values. It follows that the accelerated tests from prior work require at least a constant fraction $\geq c$ in the amount of data consumed per test compared to full-data tests, so their speed-up is $\leq 1/c$. The issue is the use of tail bounds to separate Δ from zero; for certain input/random u combinations, Δ can be arbitrarily close to zero. We avoid this by using the *approximately normal* variation in Δ^* to *replace* the variation due to u.

2.2 MCMC Posterior Inference

There is a separate line of MCMC work drawing principles from statistical physics. One can apply Hamiltonian Monte Carlo (HMC) [Neal, 2010] methods which generate high acceptance *and* distant proposals when run on full batches of data. Recently Langevin Dynamics [Welling and Teh, 2011, Ahn et al., 2012] has been applied to Bayesian estimation on minibatches of data. This simplified dynamics uses local proposals and avoids M-H tests by using small proposal steps whose acceptance approaches 1 in the limit. However, the constraint on proposal step size is severe, and the state space exploration reduces to a random walk. Full minibatch HMC for minibatches was described in Chen et al. [2014] which allows momentum-augmented proposals with larger step sizes. However, step sizes are still limited by the need to run accurately without M-H tests. By providing an M-H test with similar cost to standard gradient steps, our work opens the door to applying those methods with much more aggressive step sizes without loss of accuracy.

3 A NEW MH ACCEPTANCE TEST

3.1 Log-Likelihood Ratios

For our new M-H test, we denote the exact and approximate log likelihood ratios as Δ and Δ^*, respectively. First, Δ is defined as

$$\Delta(\theta, \theta') = \log\frac{p_0(\theta')\prod_{i=1}^{N}p(x_i \mid \theta')q(\theta \mid \theta')}{p_0(\theta)\prod_{i=1}^{N}p(x_i \mid \theta)q(\theta' \mid \theta)}, \quad (6)$$

where $p_0, p,$ and q match the corresponding functions within Equation (2). We separate out terms dependent

and independent of the data x as:

$$\Delta(\theta, \theta') = \sum_{i=1}^{N} \log \frac{p(x_i \mid \theta')}{p(x_i \mid \theta)} - \psi(1, \theta, \theta') \qquad (7)$$
$$= \Lambda(\theta, \theta') - \psi(1, \theta, \theta').$$

A minibatch estimator of Δ, denoted as Δ^*, is

$$\Delta^*(\theta, \theta') = \frac{N}{b} \sum_{i=1}^{b} \log \frac{p(x_i \mid \theta')}{p(x_i \mid \theta)} - \psi(1, \theta, \theta') \qquad (8)$$
$$= \Lambda^*(\theta, \theta') - \psi(1, \theta, \theta')$$

Note that Δ and Δ^* are evaluated on the full dataset and a minibatch of size b respectively. The term N/b means $\Delta^*(\theta, \theta')$ is an unbiased estimator of $\Delta(\theta, \theta')$.

The key to our test is a smooth acceptance function. We consider functions other than the classical Metropolis test that satisfy the detailed balance condition needed for accurate posterior estimation. A class of suitable functions is specified as follows:

Lemma 2. *If $g(s)$ is any function such that $g(s) = \exp(s)g(-s)$, then the acceptance function $\alpha(\theta, \theta') \triangleq g(\Delta(\theta, \theta'))$ satisfies detailed balance.*

This result is used in Barker [1965] to define the Barker acceptance test. As a sanity check, choosing $g(s) = \exp(s) \wedge 1$ — a function satisfying the requirement of Lemma 2 — produces the classical Metropolis acceptance test $\alpha(\theta, \theta') = g(\Delta(\theta, \theta')) = \frac{p(\theta')q(\theta|\theta')}{p(\theta)q(\theta'|\theta)} \wedge 1$. In fact, $g(s) = \exp(s) \wedge 1$ is the optimal acceptance function in terms of acceptance rate, since it accepts with probability 1 for $\Delta > 0$.

3.2 Barker (Logistic) Acceptance Function

For our new MH test we use the Barker logistic [Barker, 1965] function: $g(s) = (1 + \exp(-s))^{-1}$. Straightforward arithmetic shows that it satisfies the condition in Lemma 2. While it is slightly less efficient than the Metropolis test when used on the full dataset, we will see that its smoothness allows it to naturally tolerate substantial variance in its input. This in turn will lead to a much more efficient test on subsets of data.

Assume we begin with the current sample θ and a candidate sample θ', and that $V \sim \mathcal{U}(0, 1)$ is a uniform random variable. We accept θ' if $g(\Delta(\theta, \theta')) > V$, and reject otherwise. Since $g(s)$ is monotonically increasing, its inverse $g^{-1}(s)$ is well-defined and unique. So an equivalent test is to accept θ' iff

$$\Delta(\theta, \theta') > X = g^{-1}(V) \qquad (9)$$

where X is a random variable with the logistic distribution (its CDF is the logistic function). To see this notice

that $\frac{dV}{dX} = g'$, that g' is the density corresponding to a logistic CDF, and finally that $\frac{dV}{dX}$ is the density of X. The density of X is symmetric, so we can equivalently test whether

$$\Delta(\theta, \theta') + X > 0 \qquad (10)$$

for a logistic random variable X.

3.3 A Minibatch Acceptance Test

We now describe acceptance testing using the minibatch estimator $\Delta^*(\theta, \theta')$. From Equation (8), $\Delta^*(\theta, \theta')$ can be represented as a constant term plus the mean of b IID terms $\Lambda_i(\theta, \theta')$ of the form $N \log \frac{p(x_i|\theta')}{p(x_i|\theta)}$. As b increases, $\Delta^*(\theta, \theta')$ therefore has a distribution which approaches a normal distribution by the Central Limit Theorem. We now describe this using an asymptotic argument and defer specific bounds between the CDFs of $\Delta^*(\theta, \theta')$ and a Gaussian to Section 5.

In the limit, since Δ^* is normally distributed about its mean Δ, we can write

$$\Delta^* = \Delta + X_{\text{norm}}, \quad X_{\text{norm}} \sim \bar{\mathcal{N}}(0, \sigma^2(\Delta^*)), \qquad (11)$$

where $\bar{\mathcal{N}}(0, \sigma^2(\Delta^*))$ denotes a distribution which is approximately normal with variance $\sigma^2(\Delta^*)$. But to perform the test in Equation (10) we want $\Delta + X$ for a logistic random variable X (call it X_{log} from now on). In Bardenet et al. [2015] it was proposed to use Δ^* in a Barker test anyway and tolerate the fixed error caused by this approximation.

Our approach is to instead decompose X_{log} as

$$X_{\text{log}} = X_{\text{norm}} + X_{\text{corr}}, \qquad (12)$$

where we assume $X_{\text{norm}} \sim \mathcal{N}(0, \sigma^2)$ and that X_{corr} is a zero-mean "correction" variable with density $C_\sigma(X)$. The two variables are added (i.e., their distributions convolve) to form X_{log}. This decomposition requires an appropriate C_σ, which we derive in Section 4. Using X_{corr} samples from $C_\sigma(X)$, the acceptance test is now

$$\Delta + X_{\text{log}} = (\Delta + X_{\text{norm}}) + X_{\text{corr}} = \Delta^* + X_{\text{corr}} > 0. \qquad (13)$$

Therefore, assuming the variance of Δ^* is small enough, if we have an estimate of Δ^* from the current data minibatch, we test acceptance by adding a random variable X_{corr} and then accept θ' if the result is positive (and reject otherwise).

If $\bar{\mathcal{N}}(0, \sigma^2(\Delta^*))$ is exactly $\mathcal{N}(0, \sigma^2(\Delta^*))$, the above test is exact, and as we show in Section 5, if there is a maximum error ϵ between the CDF of $\bar{\mathcal{N}}(0, \sigma^2(\Delta^*))$ and the CDF of $\mathcal{N}(0, \sigma^2(\Delta^*))$, then our test has an error of at most ϵ relative to the full batch version.

4 CORRECTION DISTRIBUTION

Our test in Equation (13) requires knowing the distribution of X_{corr}. In Section 5, we show that the test accuracy depends on the absolute error between the CDFs of $X_{\text{norm}} + X_{\text{corr}}$ and X_{log}. Consequently, we need to minimize this in our construction of X_{corr}. More formally, let $\Phi_{s_X} = \Phi(X/s_X)$ where Φ is the standard normal CDF[3], $S(X)$ be the logistic function, and $C_\sigma(X)$ be the *density* of the correction X_{corr} distribution. Our goal is to solve:

$$C_\sigma^* = \arg\min_{C_\sigma} |\Phi_\sigma * C_\sigma - S| \qquad (14)$$

where $*$ denotes convolution.

For computation of C_σ, we assume that its input Y and another variable X lie in the intervals $[-V, V]$ and $[-2V, 2V]$, respectively. We discretize the convolution by discretizing X and Y into $4N + 1$ and $2N + 1$ values respectively. If $i \in \{-2N, \dots, 2N\} = \mathcal{I}$ and $j \in \{-N, \dots, N\} = \mathcal{J}$, then we can write $X_i = i(V/N)$ and $Y_j = j(V/N)$, and the objective can be written as:

$$C_\sigma^* = \arg\min_{C_\sigma} \max_{i \in \mathcal{I}} \left| \sum_{j \in \mathcal{J}} \Phi_\sigma(X_i - Y_j) C_\sigma(Y_j) - S(X_i) \right|.$$

We now define a matrix M and vectors u and v such that $M_{ij} = \Phi_\sigma(X_i - Y_j)$, $u_j = C_\sigma(Y_j)$, and $v_i = S(X_i)$, where the indices i and j are appropriately translated to be non-negative indices for $M, u,$ and v. Thus, the problem is now to minimize $\|Mu - v\|_\infty$ with the constraint that $u > 0$ since it represents a density. We approximate this with a least squares solution:

$$u^* = \arg\min_u \|Mu - v\|_2^2 + \lambda\|u\|_2^2, \qquad (15)$$

with regularization λ. The solution is well-known from the normal equations ($u^* = (M^T M + \lambda I)^{-1} M^T v$) and in practice yields an acceptable L_∞ norm.

With this approach, there is no guarantee that $u^* \geq 0$. However, we have some flexibility in the choice of σ in Equation (14). As we decrease the variance of X_{norm}, the variance of X_{corr} grows by the same amount and is in fact the result of convolution with a Gaussian whose variance is the difference. Thus as σ decreases, $C_\sigma(X)$ grows and approaches the derivative of a logistic function at $\sigma = 0$. It retains some weak negative values for $\sigma > 0$ but removal of those leads to small error. Table 1 shows that the errors between $X_{\text{norm}} + X_{\text{corr}}$ and X_{log} approach single floating precision (about 10^{-7}), and Algorithm 1 describes our procedure. A few points:

---

[3]Hence, Φ_{s_X} is the CDF of a zero-mean Gaussian with standard deviation s_X.

Input : Number of samples T, minibatch size m, error bound δ, pre-computed correction $C_1(X)$ distribution, initial sample θ_1.
Output : A chain of T samples $\{\theta_1, \dots, \theta_T\}$ from $p(\theta)$;
for $t = \{1, \dots, T\}$ **do**
  -Propose a candidate θ' from proposal $q(\theta' \mid \theta_t)$;
  -Draw a minibatch of m points x_i, compute $\Delta^*(\theta_t, \theta')$ and sample variance $s_{\Delta^*}^2$;
  -Estimate moments $E|\Lambda_i - \Lambda|$ and $E|\Lambda_i - \Lambda|^3$ from the sample, and error ϵ from Corollary 1;
  while $s_{\Delta^*}^2 \geq 1$ **or** $\epsilon > \delta$ **do**
    -Draw m more samples to augment the minibatch, update Δ^*, $s_{\Delta^*}^2$ and ϵ estimates;
  end
  -Draw $X_{\text{nc}} \sim \mathcal{N}(0, 1 - s_{\Delta^*}^2)$ and X_{corr} from the correction distribution $C_1(X)$;
  if $\Delta^* + X_{\text{nc}} + X_{\text{corr}} > 0$ **then**
    -Accept the candidate, $\theta_{t+1} = \theta'$;
  else
    -Reject and re-use the old sample, $\theta_{t+1} = \theta_t$;
  end
end

Algorithm 1: Our acceptance test for MCMC.

Table 1: Error (L_∞) in $X_{\text{norm}} + X_{\text{corr}}$ versus X_{log}

N	σ	λ	L_∞ error
4000	0.9	1	1.0e-4
4000	0.8	0.03	5.0e-6

- It uses an adaptive step size so as to use the smallest possible average minibatch size. Unlike previous work however (and as we show in Section 6) the size distribution is short-tailed.
- An additional normal variable X_{nc} is added to Δ^* to produce a variable with unit variance. This is not mathematically necessary, but allows us to use a single correction distribution C_1 with $\sigma = 1$ for X_{corr}, saving on memory footprint.
- The sample variance $s_{\Delta^*}^2$ is proportional to $\|\theta' - \theta\|_2^2$ whose distribution for Normal proposals is the square of a normal variable.

5 ANALYSIS

We now derive error bounds for our M-H test and the approximate target distribution it generates. In the most similar prior works, Korattikara et al. [2014] uses asymptotic arguments based on the CLT to argue that its approximate acceptance test error tends to zero as batch size increases, but no quantitative bounds are given. In Bardenet et al. [2014], explicit bounds are given, but they depend on bounding:

$$C_{\theta, \theta'} = \max_{1 \leq i \leq N} |\log p(x_i \mid \theta') - \log p(x_i \mid \theta)|. \qquad (16)$$

It is possible to analytically determine $C_{\theta,\theta'}$ for logistic regression, but it is unclear how to do so for more complicated models, and one generally needs to use all $p(x_i \mid \theta')$ terms[4] to get $C_{\theta,\theta'}$. In contrast, we use quantitative forms of the Central Limit Theorem which rely on measurable statistics from a single minibatch. Thus a sampler using our approach does not need to see data beyond the current minibatch.

In Section 5.1, we present bounds on the absolute and relative error (in terms of the CDFs) of the distribution of Δ^* vs. a Gaussian. We then show in Section 5.2 that these bounds are preserved after the addition of other random variables (e.g., X_{nc} and X_{corr}). It then follows that the acceptance test has the same error bound.

5.1 Bounding the Error of Δ^* from Gaussian

We use the following quantitative central-limit result:

Lemma 3. *Let $X_1, \ldots, X_n$ be a set of zero-mean, independent, identically-distributed random variables with sample mean $\bar{X}$ and sample variance s_X^2 where:*

$$\bar{X} = \frac{1}{n}\sum_{i=1}^{n} X_i, \quad s_X = \frac{1}{n}\left(\sum_{i=1}^{n}(X_i - \bar{X})^2\right)^{\frac{1}{2}}. \quad (17)$$

This implies $t = \bar{X}/s_X$ has an approximate Student's distribution which approaches a normal distribution in the limit. Then

$$\sup_x |\Pr(t < x) - \Phi(x)| \le \frac{6.4E|X|^3 + 2E|X|}{\sqrt{n}}. \quad (18)$$

Proof. See Supplementary Material, Section B. □

Lemma 3 demonstrates that as long as we know the first and third absolute moments $E|X|$ and $E|X|^3$, we can bound the error of the normal approximation, which decays as $O(n^{-\frac{1}{2}})$. Making the change of variables $y = x s_X$, Equation (18) becomes

$$\sup_y \left|\Pr(\bar{X} < y) - \Phi\left(\frac{y}{s_X}\right)\right| \le \frac{6.4E|X|^3 + 2E|X|}{\sqrt{n}} \quad (19)$$

showing that the distribution of $\bar{X}$ approaches the normal distribution $\mathcal{N}(0, s_X)$ whose variance is s_X, measured from the sample.

To apply this to our test, let $X_i = \Lambda_i(\theta, \theta') - \Lambda(\theta, \theta')$, so that the X_i are zero-mean, i.i.d. variables. If instead of all n samples, we only extract a subset of b samples corresponding to our minibatch, we can connect $\bar{X}$

---

[4]The sample code provided by the authors of Bardenet et al. [2014] in fact computes $C_{\theta,\theta'}$ explicitly for each sample generated, i.e. it traverses the entire data, thus providing no performance advantage over the complete test.

---

with our Δ^* term: $\bar{X} = \Delta^*(\theta, \theta') - \Delta(\theta, \theta')$, so that $s_X = s_{\Delta^*}$. This results in the following:

Corollary 1. *We can now substitute into Equation (19) and displace by the mean, giving:*

$$\sup_y \left|\Pr(\Delta^* < y) - \Phi\left(\frac{y - \Delta}{s_{\Delta^*}}\right)\right| \le \frac{6.4E|X|^3 + 2E|X|}{\sqrt{b}}$$
$$= \epsilon(\theta, \theta', b).$$

Corollary 1 shows that the distribution of Δ^* approximates a Normal distribution with mean Δ and variance $s_{\Delta^*}^2$. Furthermore, it bounds the error with *estimable quantities*: both $E|X|$ and $E|X|^3$ can be estimated as means of $|\Lambda_i - \Lambda|$ and $|\Lambda_i - \Lambda|^3$, respectively, on each minibatch. We expect this will often be accurate enough on minibatches with hundreds of points, but otherwise bootstrap CIs can be computed.

5.2 Adding Random Variables

We next relate the CDFs of distributions and show that bounds are preserved after adding random variables.

Lemma 4. *Let $P(x)$ and $Q(x)$ be two CDFs satisfying $\sup_x |P(x) - Q(x)| \le \epsilon$ with x in some real range. Let $R(y)$ be the density of another random variable y. Let P' be the convolution $P * R$ and Q' be the convolution $Q * R$. Then $P'(z)$ (resp. $Q'(z)$) is the CDF of sum $z = x + y$ of independent random variables x with CDF $P(x)$ (resp. $Q(x)$) and y with density $R(y)$. Then*

$$\sup_x |P'(x) - Q'(x)| \le \epsilon. \quad (20)$$

Proof. See Supplementary Material, Section C. □

From Lemma 4, we have the following Corollary:

Corollary 2. *If $\sup_y |\Pr(\Delta^* < y) - \Phi(\frac{y-\Delta}{s_{\Delta^*}})| \le \epsilon(\theta, \theta', b)$, then*

$$\sup_y |\Pr(\Delta^* + X_{\text{nc}} + X_{\text{corr}} < y) - S(y - \Delta)| \le \epsilon(\theta, \theta', b)$$

where $S(x)$ is the standard logistic function, and X_{nc} and X_{corr} are generated as per Algorithm 1.

Proof. See Supplementary Material, Section D. □

From Section 3, as the distribution of Δ^* approaches a Gaussian, our new MH test becomes more accurate. Corollary 2 shows that the bounds from Section 5.1 are preserved after the addition of the random variables we use, showing that our test should remain accurate. In fact we can do better ($O(n^{-1})$) by using a more precise limit distribution under an additional assumption. We review this in the Supplementary Material, Section E, and leave the details to future work.

Table 2: Gaussian Mixture Model Statistics

Metric	Ours	Korat.'14	Barde.'14
Equation 22	-1430.0	-1578.9	-1232.7
Chi-Squared	3313.9	3647.7	2444.1

6 EXPERIMENTS

We conduct two sets of experiments analyzing convergence to the posterior of a heated Gaussian mixture, and efficiency on logistic regression.

6.1 Mixture of Gaussians

We borrow a Gaussian mixture experiment from Welling and Teh [2011]. The parameter is 2-D, $\theta = \langle \theta_1, \theta_2 \rangle$, and the generation process is

$$
\begin{aligned}
\theta &\sim \mathcal{N}(0, \mathrm{diag}(\sigma_1^2, \sigma_2^2)) \\
x_i &\sim 0.5 \cdot \mathcal{N}(\theta_1, \sigma_x^2) + 0.5 \cdot \mathcal{N}(\theta_1 + \theta_2, \sigma_x^2).
\end{aligned}
\tag{21}
$$

We set $\sigma_1^2 = 10, \sigma_2^2 = 1$ and $\sigma_x^2 = 2$. We fix $\theta = \langle 0, 1 \rangle$. The original paper sampled 100 data points and estimated the posterior. We are interested in performance on larger problems and so sampled 1,000,000 points to form the posterior of $p(\theta) \prod_{i=1}^{1,000,000} p(x_i \mid \theta)$ with the same prior from Equation (21). This produces a much sharper posterior with two very narrow peaks. Our goal is to reproduce the original posterior, so we adjust the temperature to $T = 10,000$. Taking logs, we get the target as shown in the far left of Figure 1.

We run MCMC with our test and benchmark with Korattikara et al. [2014] and Bardenet et al. [2014], all with initial minibatch size 100. For the former, we increment sizes by 100 and set the tolerance $\epsilon = 0.005$. For the latter, we increase sizes geometrically with $\gamma = 1.5$ and use parameters $p = 2, \delta = 0.01$. All methods collect 5000 samples using a random walk proposer with covariance $\mathrm{diag}(0.15, 0.15)$, which means the M-H test is responsible for the sample distribution.

Figure 1 shows scatter plots of the resulting θ samples for the three methods, with darker regions indicating a greater density of points. There are no obvious differences, so we measure the similarity between each set of samples and the actual posterior.

We discretize the posterior coordinates into bins with respect to the two components of θ. The probability P_i of a sample falling into bin i is the integral of the true posterior over the bin's area. A single sample should therefore be multinomial with distribution P, and a set of n (ideally independent) samples is Multinomial(P, n). This distribution is simple and we can use it to measure the quality of the samples rather than use general purpose tests like KL-divergence or likelihood-ratio,

which can be problematic with zero counts in bins.

For large n, the per-bin distributions are approximated by Poissons with parameter $\lambda_i = P_i n$. Given samples $\{\theta_1, \ldots, \theta_T\}$, let c_j denote the number of individual samples θ_i that fall in bin j out of N_{bins} total. We have

$$
\log p(c_1, \ldots, c_{N_{\mathrm{bins}}} \mid P_1, \ldots, P_{N_{\mathrm{bins}}}) =
$$
$$
\sum_{j=1}^{N_{\mathrm{bins}}} c_j \log(nP_j) - nP_j - \log(\Gamma(c_j + 1)).
\tag{22}
$$

Table 2 shows the results, but it is difficult to interpret the scores, so we perform significance tests to show the difference between the MCMC-samples and ground-truth using the Chi-Squared distribution as the test statistic (also in Table 2). Both metrics imply that our method is slightly superior to Korattikara et al. [2014], but slightly worse than Bardenet et al. [2014].

Figure 2, however, suggests that our method dominates in terms of speed and efficiency. The histograms of the (final) minibatch sizes used each iteration show that our method consumes significantly less data; the distribution is short-tailed and the mean is 210. The other methods occasionally need to use nearly all data points, and average minibatch sizes are 15562 and 16857. The sizes roughly predict the running times of these methods, which are proportional to the data consumed, with the exception of Bardenet et al. [2014], which requires a full pass to compute $C_{\theta, \theta'}$ (Equation (16)). Using minibatch sizes as a proxy for time also avoids discrepancies due to code optimization.

6.2 Logistic Regression

We next use logistic regression for the binary classification of 1s versus 7s on a subset of the MNIST8M dataset, a larger version of MNIST [LeCun and Cortes] and well-suited to our big data objectives. We randomly subsampled 450k training and 192k testing points. The pixels are scaled in $[0, 1]$. We impose a flat prior on θ and again use a random walk proposer, this time with covariance matrix $0.05I$ for the 784×784 identity matrix I. The temperature is set at $T = 1000$. We run the three tests for samples and again use 100 as the starting minibatch size. To make Korattikara et al. [2014] run faster, we now use $\epsilon = 0.05$. For Bardenet et al. [2014], we tuned $\gamma \in \{1.1, 1.25, 1.5, 1.75, 2\}$ before concluding that $\gamma = 2.0$ was OK. In addition, we tried their analytic approximation to $C_{\theta, \theta'}$, but found that the bounds were too high, meaning that their method needed to consume the entire dataset each iteration. Thus, we naively compute $C_{\theta, \theta'}$.

Figure 3 shows the test log likelihood and prediction accuracy as a function of the cumulative training points

Figure 1: The log posterior contours and scatter plots of sampled θ values using different methods.


Figure 2: Minibatch sizes used in Section 6.1's experiment. The axes have the same (log-log scale) range.

processed.[5] To generate the curves, for each of the sampled vectors θ_t, $t \in \{1, \ldots, 3000\}$, we use θ_t as the parameter for logistic regression. Our minibatch MH test is more efficient in terms of accuracy, achieving convergence using roughly half as many elements compared to Korattikara et al. [2014] and **TODO TODO TODO** compared to Bardenet et al. [2014], though its log likelihood results lag slightly behind the former algorithm during the very early stages.

Figure 4 shows histograms of minibatch sizes for the three methods on a log-log scale. With an initial size of 100, we achieve an average minibatch size of 393, far smaller than the averages of the other two methods, which have longer-tailed distributions.

7 CONCLUSIONS

In this paper, we derived an M-H test for minibatch MCMC which approximates full data tests. We provide theoretical results and experimentally show the benefits of our test on Gaussian mixtures and logistic regression. Directions for future work include running more experiments with a particular focus on controlling variances. More elaborate extensions include testing on neural networks, combining our results with Hamiltonian Monte Carlo methods, providing a recipe for how to use our algorithm (following the framework of Ma et al. [2015]), or integrating parallel MCMC [Angelino et al., 2014, Ahn et al., 2014] concepts.

---

[5]Note that the curves do not span the same length over the x-axis, because our test consumes fewer samples throughout the MCMC procedure, so the corresponding curve will "end" before the other two.

References

S. Ahn, A. K. Balan, and M. Welling. Bayesian posterior sampling via stochastic gradient fisher scoring. In *Proceedings of the 29th International Conference on Machine Learning (ICML)*, 2012.

S. Ahn, B. Shahbaba, and M. Welling. Distributed stochastic gradient MCMC. In *Proceedings of the 31st International Conference on Machine Learning, (ICML)*, 2014.

E. Angelino, E. Kohler, A. Waterland, M. Seltzer, and R. P. Adams. Accelerating MCMC via parallel predictive prefetching. In *Proceedings of the Thirtieth Conference on Uncertainty in Artificial Intelligence, (UAI)*, 2014.

R. Bardenet, A. Doucet, and C. Holmes. Towards scaling up markov chain monte carlo: an adaptive subsampling approach. In *Proceedings of the 31st International Conference on Machine Learning (ICML)*, 2014.

R. Bardenet, A. Doucet, and C. Holmes. On markov chain monte carlo methods for tall data. *arXiv preprint arXiv:1505.02827*, 2015.

A. A. Barker. Monte-carlo calculations of the radial distribution functions for a proton-electron plasma. *Australian Journal of Physics*, 18:119–133, 1965.

V. Bentkus, F. Gotze, and W.R.vanZwet. An edgeworth expansion for symmetric statistics. *Annals of Statistics*, 25(2), 1997.

Figure 3: Logistic regression performance (accuracy/log likelihood) based on cumulative data usage.


Figure 4: Counts of minibatch sizes in the MNIST logistic regression experiment (analogous to Figure 2).

S. Brooks, A. Gelman, G. Jones, and X.-L. Meng. *Handbook of Markov Chain Monte Carlo*. CRC press, 2011.

T. Chen, E. Fox, and C. Guestrin. Stochastic gradient Hamiltonian Monte Carlo. In *Proceedings of the 31st International Conference on Machine Learning (ICML)*, 2014.

A. Choromanska, M. Henaff, M. Mathieu, G. B. Arous, and Y. LeCun. The loss surface of multilayer networks. In *Proceedings of the 18th International Conference on Artificial Intelligence and Statistics, AISTATS*, 2015.

W. Gilks and D. Spiegelhalter. *Markov chain Monte Carlo in practice*. Chapman & Hall/CRC, 1996.

W. K. Hastings. Monte carlo sampling methods using markov chains and their applications. *Biometrika*, 57: 97–109, 1970.

A. Korattikara, Y. Chen, and M. Welling. Austerity in MCMC land: Cutting the metropolis-hastings budget. In *Proceedings of the 31st International Conference on Machine Learning (ICML)*, 2014.

Y. LeCun and C. Cortes. MNIST handwritten digit database. URL http://yann.lecun.com/exdb/mnist/.

Y. Ma, T. Chen, and E. Fox. A complete recipe for stochastic gradient mcmc. In *Advances in Neural Information Processing Systems 28*, 2015.

D. Maclaurin and R. P. Adams. Firefly monte carlo: Exact MCMC with subsets of data. In *Proceedings of the Thirtieth Conference on Uncertainty in Artificial Intelligence, (UAI)*, 2014.

N. Metropolis, A. W. Rosenbluth, M. N. Rosenbluth, A. H. Teller, and E. Teller. Equation of state calculations by fast computing machines. *The Journal of Chemical Physics*, 21, 1953.

R. M. Neal. MCMC using Hamiltonian dynamics. *Handbook of Markov Chain Monte Carlo*, 54:113–162, 2010.

Y. Novak. On self-normalized sums and students statistic. *Theory of Probability and its Applications*, 49(2): 336–344, 2005.

G. O. Roberts and J. S. Rosenthal. Optimal scaling for various metropolishastings algorithms. *Statistical Science*, 16(4):351367, 2001.

M. Welling and Y. W. Teh. Bayesian learning via stochastic gradient langevin dynamics. In *Proceedings of the 28th International Conference on Machine Learning (ICML)*, 2011.

Supplementary Material

A Proof of Lemma 1

Choose $(\theta' - \theta) \in \pm\frac{1}{\sqrt{N}}[0.5, 1]$ (event 1) and $(\theta - 0.5) \in \pm\frac{1}{\sqrt{N}}[0.5, 1]$ filtered for matching sign (event 2). As discussed in Lemma 1, both $q(\theta' \mid \theta)$ and $p(\theta \mid x_1, \ldots, x_N)$ have variance $1/N$. If we denote Φ as the CDF of the standard normal distribution, then the former event occurs with probability $p_0 = 2(\Phi(\sqrt{N}\frac{1}{\sqrt{N}}) - \Phi(\sqrt{N}\frac{0.5}{\sqrt{N}})) = 2(\Phi(1) - \Phi(0.5)) \approx 0.2997$. The latter event, because we restrict signs, occurs with probability $p_1 = \Phi(1) - \Phi(0.5) \approx 0.14988$.

These events together guarantee that $\Lambda^*(\theta, \theta')$ is negative by inspection of equation (24) below. This implies that we can find a $u \in (0, 1)$ so that $\psi(u, \theta, \theta') = \log u < 0$ equals $E[\Lambda^*(\theta, \theta')]$. Specifically, choose u_0 to satisfy $\log u_0 = E[\Lambda^*(\theta, \theta')]$. Using $E[x_i] = 0.5$ and Equation (5), we see that

$$\log u_0 = N(\theta' - \theta)\frac{1}{b} \cdot E\left[\sum_{i=1}^{b} x_i - \theta - \frac{\theta' - \theta}{2}\right] \quad (23)$$

$$\log u_0 = -N(\theta' - \theta)\left(\theta - 0.5 + \frac{\theta' - \theta}{2}\right). \quad (24)$$

Next, consider the minibatch acceptance test $\Lambda^*(\theta, \theta') \not\approx \psi(u, \theta, \theta')$ used in Korattikara et al. [2014] and Bardenet et al. [2014], where $\not\approx$ means "significantly different from" under the distribution over samples of x_i. This turns out to be

$$\Lambda^*(\theta, \theta') \not\approx \psi(u_0, \theta, \theta')$$

$$\iff N(\theta' - \theta) \cdot \frac{1}{b}\sum_{i=1}^{b} x_i - \theta - \frac{\theta' - \theta}{2} \not\approx \log u_0$$

$$\iff \frac{1}{b}\sum_{i=1}^{b} x_i - \left(\theta + \frac{\theta' - \theta}{2} + \frac{\log u_0}{N(\theta' - \theta)}\right) \not\approx 0$$

$$\iff \frac{1}{b}\sum_{i=1}^{b} x_i - 0.5 \not\approx 0.$$

$$(25)$$

Since the x_i have mean 0.5, the resulting test with our chosen u_0 will never correctly succeed and must use all N data points. Furthermore, if we sample values of u near enough to u_0, the terms in parenthesis will not be sufficiently different from 0.5 to allow the test to succeed.

The choices above for θ and θ' guarantee that

$$\log u_0 \in -[0.5, 1][0.75, 1.5] = [-1.5, -0.375]. \quad (26)$$

Next, consider the range of u values near u_0:

$$\log u \in \log u_0 + [-0.5, 0.375]. \quad (27)$$

The size of the range in u is at least $\exp([-2, -1.125]) \approx [0.13534, 0.32465]$ and occurs with probability at least $p_2 = 0.18932$. With u in this range, we rewrite the test as:

$$\frac{1}{b}\sum_{i=1}^{b} x_i - 0.5 \quad \not\approx \quad \frac{\log u/u_0}{N(\theta' - \theta)} \quad (28)$$

so that, as in Equation (25), the LHS has expected value zero. Given our choice of intervals for the variables, we can compute the range for the right hand side (RHS) assuming[6] that $\theta' - \theta > 0$:

$$\min\{\text{RHS}\} = \frac{-0.5}{\sqrt{N} \cdot 0.5} = -\frac{1}{\sqrt{N}}$$

$$\text{and} \quad \max\{\text{RHS}\} = \frac{0.375}{\sqrt{N} \cdot 0.5} = \frac{0.75}{\sqrt{N}} \quad (29)$$

Thus, the RHS is in $\frac{1}{\sqrt{N}}[-1, 0.75]$. The standard deviation of the LHS given the interval constraints is at least $0.5/\sqrt{b}$. Consequently, the gap between the LHS and RHS in Equation (28) is at most $2\sqrt{b/N}$ standard deviations, limiting the range in which the test will be able to "succeed" without requiring more samples.

The samples θ, θ' and u are drawn independently and so the probability of the conjunction of these events is $c = p_0 p_1 p_2 = 0.0085$.

B Proof of Lemma 3

The following bound is given immediately after Corollary 2 from Novak [2005]:

$$-6.4E|X|^3 - 2E|X| \leq \sup_x |\Pr(t < x) - \Phi(x)|\sqrt{n}$$

$$\leq 1.36E|X|^3.$$

This bound applies to $x \geq 0$. Applying the bound to $-x$ when $x < 0$ and combining with $x > 0$, we obtain the weaker but unqualified bound in Equation (18).

C Proof of Lemma 4

We first observe that

$$P'(z) - Q'(z) = \int_{-\infty}^{+\infty} (P(z - x) - Q(z - x))R(x)dx,$$

and since $\sup_x |P(x) - Q(x)| \leq \epsilon$ it follows that $\forall z$:

$$-\epsilon = \int_{-\infty}^{+\infty} -\epsilon R(x)dx$$

$$\leq \int_{-\infty}^{+\infty} (P(z - x) - Q(z - x))R(x)dx$$

$$\leq \int_{-\infty}^{+\infty} \epsilon R(x)dx = \epsilon,$$

---

[6]If $\theta' - \theta < 0$, then the range would be $\frac{1}{\sqrt{N}}[-0.75, 1]$ but this does not matter for the purposes of our analysis.

as desired.

D Proof of Corollary 2

We apply Lemma 4 twice. First take:

$$P(y) = \Pr(\Delta^* < y)$$
$$\text{and} \quad Q(y) = \Phi\left(\frac{y - \Delta}{s_{\Delta^*}}\right) \tag{30}$$

and convolve with the distribution of X_n which has density $\phi(X/\sigma_n)$ where $\sigma_n^2 = 1 - s_{\Delta^*}^2$. This yields the next iteration of P and Q:

$$P'(y) = \Pr(\Delta^* + X_{\text{nc}} < y)$$
$$\text{and} \quad Q'(y) = \Phi(y - \Delta) \tag{31}$$

Now we convolve with the distribution of X_{corr}:

$$P''(y) = \Pr(\Delta^* + X_{\text{nc}} + X_{\text{corr}} < y)$$
$$\text{and} \quad Q''(y) = S(y - \Delta) \tag{32}$$

Both steps preserve the error bound $\epsilon(\theta, \theta', b)$. Finally $S(y - \Delta)$ is a logistic CDF centered at Δ, and so $S(y - \Delta) = \Pr(\Delta + X_{\log} < y)$ for a logistic random $X_{\log}$. We conclude that the probability of acceptance for the actual test $\Pr(\Delta^* + X_{\text{nc}} + X_{\text{corr}} > 0)$ differs from the exact test $\Pr(\Delta + X_{\log} > 0)$ by at most ϵ.

E Improved Error Bounds Based on Skew Estimation

We demonstrate how we can show $O(n^{-1})$ error for the quantitative CLT by using a more precise limit distribution under an additional assumption. Let μ_i denote the i^{th} moment, and b_i denote the i^{th} absolute moment of X. If Cramer's condition holds:

$$\lim_{t \to \infty} \sup |E(\exp(itX))| < 1, \tag{33}$$

then Equation 2.2 in Bentkus et al.'s work on Edgeworth expansions [Bentkus et al., 1997] provides:

Lemma 5. *Let $X_1, \ldots, X_n$ be a set of zero-mean, independent, identically-distributed random variables with sample mean $\hat{X}$ and with t defined as in Lemma 3. If X satisfies Cramer's condition, then*

$$\sup_x \left| \Pr(t < x) - G\left(x, \frac{\mu_3}{b_2^{3/2}}\right) \right| \leq \frac{c(\epsilon, b_2, b_3, b_4, b_{4+\epsilon})}{n}$$

where

$$G_n(x, y) = \Phi(x) + \frac{y(2x^2 + 1)}{6\sqrt{n}} \Phi'(x). \tag{34}$$

Lemma 5 shows that the average of the X_i has a more precise, skewed CDF limit $G_n(x, y)$ where the skew term has weight proportional to a certain measure of skew derived from the moments: $\mu_3/b_2^{3/2}$. Note that if the X_i are symmetric, the weight of the correction term is zero, and the CDF of the average of the X_i converges to $\Phi(x)$ at a rate of $O(n^{-1})$.

Here the limit $G_n(x, y)$ is a normal CDF plus a correction term that decays as $n^{-1/2}$. Importantly, since $\phi''(x) = x^2\phi(x) - \phi(x)$ where $\phi(x) = \Phi'(x)$, the correction term can be rewritten giving:

$$G_n(x, y) = \Phi(x) + \frac{y}{6\sqrt{n}}(2\phi''(x) + 3\phi(x)) \tag{35}$$

From which we see that $G_n(x, y)$ is a linear combination of $\Phi(x)$, $\phi(x)$ and $\phi''(x)$. In Algorithm 1, we correct for the difference in σ between Δ^* and the variance needed by X_{corr} using X_{nc}. This same method works when we wish to estimate the error in Δ^* vs $G_n(x, y)$. Since all of the component functions of $G_n(x, y)$ are derivatives of a (unit variance) $\Phi(x)$, adding a normal variable with variance σ' increases the variance of all three functions to $1 + \sigma'$. Thus we add X_{nc} as per Algorithm 1 preserving the limit in Equation (35).

The deconvolution approach can be used to construct a correction variable X_{corr} between $G_n(x, y)$ and $S(x)$ the standard logistic function. An additional complexity is that $G_n(x, y)$ has additional parameters y and n. Since these act as a single multiplier $\frac{y}{6\sqrt{n}}$ in Equation (35), its enough to consider a function $g(x, y')$ parametrized by $y' = \frac{y}{6\sqrt{n}}$. This function can be computed and saved offline. As we have shown earlier, errors in the "limit" function propagate directly through as errors in the acceptance test. To achieve a test error of 10^{-6} (close to single floating point precision), we need a y' spacing of 10^{-6}. It should not be necessary to tabulate values all the way to $y' = 1$, since y' is scaled inversely by the square root of minibatch size. Assuming a max y' of 0.1 requires us to tabulate about 100,000. Since our x resolution is 10,000, this leads to a table with about 1 billion values, which can comfortably be stored in memory. However, if $g(x, y)$ is moderately smooth in y, it should be possible to achieve similar accuracy with a much smaller table. We leave further analysis and experiments with $g(x, y)$ as future work.

F NIPS 2016 Submission Statement

We previously submitted a much older draft of this paper to NIPS 2016. The current manuscript has been substantially revised since that submission. We have made the following changes:

1. At the time of the NIPS submission, we were unaware of several references which described some of what we wrote. For this submission, we now cite the important work of [Bardenet et al., 2015] which had the idea of using subsampling noise for exploration (rather than the $\log u$ variable). We also cite the work of [Barker, 1965] which uses the Barker function. We welcome information about other potentially missing references.

2. The previous theoretical results have been entirely scrapped and replaced with different but more relevant results. The NIPS submission listed a page of theoretical results which did not tie into the rest of the paper's contributions and were confusing to the readers. We now have theoretical results that are much clearer and also *directly* show the expected performance benefits of our acceptance test. Note also that Lemma 1 is an entirely new addition to this version.

3. Several reviewers mentioned that our deconvolution approach to determine the correction distribution was not unique and thus ill-defined. Consequently, we have added an entirely new section of the paper (Section 4) explaining how we derived that distribution.

4. Reviewers pointed out one of our mistakes when we said our minibatch sizes were fixed. We have changed this in the current manuscript to mean that the per-iteration minibatch sizes for our test have a distribution with *shorter tails* than those of prior work, which can be observed by plotting a histogram of minibatch sizes.

5. Our experiments now include new comparisons with a baseline from [Bardenet et al., 2014] (in addition to [Korattikara et al., 2014], which we had earlier). In addition, each algorithm now runs MCMC sampling on the *same* distribution. Previously, we ran our distribution at a higher temperature but kept the algorithm from [Korattikara et al., 2014] running on the distribution at temperature $T = 1$. In this set of experiments, we tuned hyperparameters of the other algorithms to make the comparison fairer. Finally, in the Gaussian mixture model scenario (Section 6.1), we provide more details on how "accurate" the samples are, rather than solely relying on the visualization of Figure 1.

6. We improved Algorithm 1 so that we show explicitly when we compute the moments, and why we now only need one distribution C_σ for $\sigma = 1$ due to the extra X_{nc} variable.

7. Finally, we made minor revisions addressing: differences between the Barker function vs. the original MH test, and sampling with vs. without replacement.

We are confident that the current manuscript is of far better quality than the NIPS submission, and we appreciate the efforts of the NIPS reviewers to give us ideas for improvement.