
An Efficient Minibatch Acceptance Test for Metropolis-Hastings

Anonymous Author 1
Unknown Institution 1

Anonymous Author 2
Unknown Institution 2

Anonymous Author 3
Unknown Institution 3

Abstract

We present a novel Metropolis-Hastings method for large datasets that uses small expected-size minibatches of data. Previous work on reducing the cost of Metropolis-Hastings tests yield variable data consumed per sample, with only constant factor reductions versus using the full dataset for each sample. Here we present a method that can be tuned to provide arbitrarily small batch sizes, by adjusting either proposal step size or temperature. Our test uses the noise-tolerant Barker acceptance test with a novel additive correction variable. The resulting test has similar cost to a normal SGD update. Our experiments demonstrate several order-of-magnitude speedups over previous work.

1 INTRODUCTION

Markov chain Monte Carlo (MCMC) sampling is a powerful method for computation on intractable distributions. We are interested in large dataset applications, where the goal is to sample a posterior distribution $p(\theta \mid x_1, \dots, x_N)$ of parameter θ for large N . The Metropolis-Hastings method (M-H) generates sample candidates from a proposal distribution q which is in general different from the target distribution p , and decides whether to accept or reject based on an acceptance test. The acceptance test is usually a Metropolis test [Metropolis et al., 1953, Hastings, 1970].

Many state-of-the-art machine learning methods, and deep learning in particular, are based on minibatch updates (such as SGD) to a model. Minibatch updates produce many improvements to the model for each pass over the dataset, and have high sample efficiency. In contrast, conventional M-H requires calculations over

the full dataset to produce a new sample. Recent results from Korattikara et al. [2014] and Bardenet et al. [2014] perform approximate (bounded error) acceptance tests using subsets of the full dataset. The amount of data consumed for each test varies significantly from one minibatch to the next. By contrast, Maclaurin and Adams [2014], Bardenet et al. [2015] perform exact tests but require a lower bound on the parameter distribution across its domain. The amount of data reduction depends on the accuracy of this bound, and such bounds are only available for relatively simple distributions.

Here we derive a new test which incorporates the variability in minibatch statistics as *a natural part of the test* and requires less data per iteration than prior work. We use a Barker test function [Barker, 1965], which makes our test naturally error tolerant. The idea of using a noise-tolerant Barker’s test function was suggested but not explored empirically in Bardenet et al. [2015] section 6.3. But the asymptotic test statistic CDF and the Barker function are different, which leads to fixed errors for the approach in Bardenet et al. [2015]. Here, we show that the difference between the distributions can be corrected with an additive random variable. This leads to a test which is fast, and whose error can be made arbitrarily small.

Our test is applicable when the variance (over data samples) of the log acceptance probability is small enough (less than 1). It’s not clear at first why this quantity should be bounded, but we will show that it is “natural” for well-specified models running Metropolis-Hastings sampling with optimal proposals [Roberts and Rosenthal, 2001] on a full dataset. When we reduce the amount of data for the test, the variance goes up. We have to reduce variance in one of several ways. Either:

- Increase the temperature of the target distribution. Log likelihoods scale as $1/T$, and so the variance of the likelihood ratio will vary as $1/T^2$. Our model is no longer well-specified (we are doing inference at a temperature different from that assumed during data generation), but higher temperature can be advantageous for parameter exploration.

- For continuous distributions, reduce the proposal step size and variance compared to an optimal proposal. The variance of the log acceptance probability scales as the square of proposal step size.

It is worth discussing at this point the typical goals of M-H sampling on large datasets. By the Bernstein-von Mises Theorem, the posterior distribution for a Bayesian inference task has variance that scales inversely with N . Simply sampling from it is one application, but an efficient proposal [Roberts and Rosenthal, 2001] has similar variance to the target and will diffuse to it extremely slowly. For applications to neural networks or models where the posterior is multimodal [Choromanska et al., 2015], samplers will likely get trapped in one of the modes. A common solution is to anneal the sampler, running first at high temperatures to flatten the likelihood landscape. This in turn reduces the variance of the log acceptance probability and allows our test to be applied. These samples can cover the search space densely with small steps rather than taking a few sparse steps towards an optimum. In this mode, Metropolis-Hastings can be used in similar fashion to Stochastic Gradient Descent. The goal in SGD is to make gradual progress to a posterior mode with each step, taking small steps so that the cumulative displacement has progressively lower variance.

The contributions of this paper are as follows:

- We develop a new, more efficient (in samples per test) minibatch acceptance test with quantifiable error bounds. The test uses a novel additive correction variable to implement a Barker test based on minibatch mean and variance.
- We compare performance of our new test and prior approaches on several datasets. We demonstrate several order-of-magnitude improvements in efficiency (measured as data consumed per test), and show that it does not suffer from long-tailed minibatch sizes.

2 PRELIMINARIES

In the Metropolis-Hastings method [Gills and Spiegelhalter, 1996, Brooks et al., 2011], a difficult-to-compute probability distribution $p(\theta)$ is sampled using a Markov chain $\theta_1, \dots, \theta_T$. The sample θ_{t+1} at time $t + 1$ is generated using a candidate θ' from a (simpler) proposal distribution $q(\theta' | \theta_t)$, filtered by an acceptance test. The acceptance test is usually a Metropolis test. The Metropolis test has acceptance probability:

$$\alpha(\theta_t, \theta') = \frac{p(\theta')q(\theta_t | \theta')}{p(\theta_t)q(\theta' | \theta_t)} \wedge 1 \quad (1)$$

where $a \wedge b$ denotes $\min(a, b)$. With probability $\alpha(\theta_t, \theta')$, we accept θ' and set $\theta_{t+1} = \theta'$, otherwise set $\theta_{t+1} = \theta_t$. The test is often implemented with an auxiliary random variable $u \sim \mathcal{U}(0, 1)$ with a comparison $u < \alpha(\theta_t, \theta')$; here, $\mathcal{U}(a, b)$ denotes the uniform distribution on the interval $[a, b]$. For simplicity, we drop the subscript t for the current sample θ_t and denote it as θ .

The acceptance test guarantees detailed balance, which means $p(\theta)p(\theta' | \theta) = p(\theta')p(\theta | \theta')$, where $p(\theta' | \theta)$ is the probability of a transition from state θ to θ' . Here, $p(\theta' | \theta) = q(\theta' | \theta)\alpha(\theta, \theta')$. This condition, together with ergodicity, guarantees that the Markov chain has a unique stationary distribution $\pi(\theta) = p(\theta)$. For Bayesian inference, the target distribution is $p(\theta | x_1, \dots, x_N)$. The acceptance probability is now:

$$\alpha(\theta, \theta') = \frac{p_0(\theta') \prod_{i=1}^N p(x_i | \theta') q(\theta | \theta')}{p_0(\theta) \prod_{i=1}^N p(x_i | \theta) q(\theta' | \theta)} \wedge 1 \quad (2)$$

where $p_0(\theta)$ is the prior. Computing samples this way requires all N data points, but this is very expensive for large datasets. To address this challenge, Korattikara et al. [2014], Bardenet et al. [2014] perform approximate Metropolis-Hastings tests using sequential hypothesis testing. Each iteration, they start with a small minibatch and test whether θ' should be accepted based on approximating $u < \alpha(\theta, \theta')$. If the approximate test cannot decide with sufficient confidence, the minibatch size is increased and the test repeats. This process continues until a decision. The bounds depend on either an asymptotic Central Limit Theorem [Korattikara et al., 2014] or a concentration bound [Bardenet et al., 2014]. The latter requires direct bounds on the log likelihood ratio, which for general distributions requires knowing $p(x_i | \theta)$ and $p(x_i | \theta')$ for all N samples. In addition, while both methods show useful reductions in the number of samples required, they suffer the drawback of resolving small log likelihood ratio differences between the minibatch and full batch versions. We discuss a worst-case scenario in Section 2.1.

Following Bardenet et al. [2014], we write the test $u < \alpha(\theta, \theta')$ equivalently as $\Lambda(\theta, \theta') > \psi(u, \theta, \theta')$, where¹

$$\Lambda(\theta, \theta') = \sum_{i=1}^N \log \left(\frac{p(x_i | \theta')}{p(x_i | \theta)} \right) \quad (3)$$

and $\psi(u, \theta, \theta') = \log \left(u \frac{q(\theta' | \theta) p_0(\theta)}{q(\theta | \theta') p_0(\theta')} \right)$.

To reduce computational effort, an unbiased estimate of $\Lambda(\theta, \theta')$ based on a minibatch can be used:

$$\Lambda^*(\theta, \theta') = \frac{N}{b} \sum_{i=1}^b \log \left(\frac{p(x_i | \theta')}{p(x_i | \theta)} \right) \quad (4)$$

¹Our definitions differ from those in Bardenet et al. [2014] by a factor of N to simplify our analysis later.

Finally, it will be convenient for our analysis to define $\Lambda_i(\theta, \theta') = N \log(\frac{p(x_i|\theta')}{p(x_i|\theta)})$. Thus, $\Lambda(\theta, \theta')$ is the mean of $\Lambda_i(\theta, \theta')$ over the entire dataset, and $\Lambda^*(\theta, \theta')$ is the mean of $\Lambda_i(\theta, \theta')$ over its minibatch.

Since minibatches contains randomly selected samples x_i , the values Λ_i are i.i.d. random variables.² By the Central Limit Theorem, we expect $\Lambda^*(\theta, \theta')$ to be approximately Gaussian. The acceptance test then becomes a statistical test of the hypothesis that $\Lambda(\theta, \theta') > \psi(u, \theta, \theta')$ by establishing that $\Lambda^*(\theta, \theta')$ is substantially larger than $\psi(u, \theta, \theta')$.

2.1 A Worst-Case Gaussian Example

Let x_1, \dots, x_N be i.i.d. $\mathcal{N}(\theta, 1)$ with known variance $\sigma^2 = 1$ and (unknown) mean $\theta = 0.5$. We use a uniform prior on θ . The log likelihood ratio is

$$\Lambda^*(\theta, \theta') = N(\theta' - \theta) \left(\frac{1}{b} \sum_{i=1}^b x_i - \theta - \frac{\theta' - \theta}{2} \right) \quad (5)$$

which is normally distributed over selection of the Normal samples x_i . Since the x_i have unit variance, their mean has variance $1/b$, and the variance of $\Lambda^*(\theta, \theta')$ is $\sigma^2(\Lambda^*) = (\theta' - \theta)^2 N^2/b$. In order to pass a hypothesis test that $\Lambda > \psi$, there needs to be a large enough gap (several $\sigma(\Lambda^*)$) between $\Lambda^*(\theta, \theta')$ and $\psi(u, \theta, \theta')$.

The posterior is a Gaussian centered on the sample mean μ , and with variance $1/N$ (i.e., $\mathcal{N}(\mu, 1/N)$). In one dimension, an efficient proposal distribution has the same variance as the target distribution [Roberts and Rosenthal, 2001], so we use a proposal based on $\mathcal{N}(\theta, 1/N)$. It is symmetric $q(\theta' | \theta) = q(\theta | \theta')$, and since we assumed a uniform prior, $\psi(u, \theta, \theta') = \log u$. Our worst-case scenario is specified in Lemma 1.

Lemma 1. *For the model in Section 2.1, there exists a fixed (independent of N) constant c such that with probability $\geq c$ over the joint distribution of (θ, θ', u) , the tests from Korattikara et al. [2014], Bardenet et al. [2014] consume all N samples.*

Proof. See Supplementary Material, Section A. \square

Similar results can be shown for other distributions and proposals by identifying regions in product space (θ, θ', u) such that the hypothesis test needs to separate nearly-equal values. It follows that the accelerated tests from prior work require at least a constant fraction $\geq c$ in the amount of data consumed per test compared to full-data tests, so their speed-up is $\leq 1/c$. The issue is the use of tail bounds to separate $\Lambda - \psi$ from zero; for

certain input/random u combinations, this difference can be arbitrarily close to zero. We avoid this by using the *approximately normal* variation in Λ^* to *replace* the variation due to u .

2.2 MCMC Posterior Inference

There is a separate line of MCMC work drawing principles from statistical physics. One can apply Hamiltonian Monte Carlo (HMC) [Neal, 2010] methods which generate high acceptance *and* distant proposals when run on full batches of data. Recently Langevin Dynamics [Welling and Teh, 2011, Ahn et al., 2012] has been applied to Bayesian estimation on minibatches of data. This simplified dynamics uses local proposals and avoids M-H tests by using small proposal steps whose acceptance approaches 1 in the limit. However, the constraint on proposal step size is severe, and the state space exploration reduces to a random walk. Full minibatch HMC for minibatches was described in Chen et al. [2014] which allows momentum-augmented proposals with larger step sizes. However, step sizes are still limited by the need to run accurately without M-H tests. By providing an M-H test with similar cost to standard gradient steps, our work opens the door to applying those methods with much more aggressive step sizes without loss of accuracy.

3 A NEW MH ACCEPTANCE TEST

3.1 Log-Likelihood Ratios

For our new M-H test, we denote the exact and approximate log likelihood ratios as Δ and Δ^* , respectively. First, Δ is defined as

$$\Delta(\theta, \theta') = \log \frac{p_0(\theta') \prod_{i=1}^N p(x_i | \theta') q(\theta | \theta')}{p_0(\theta) \prod_{i=1}^N p(x_i | \theta) q(\theta' | \theta)}, \quad (6)$$

where p_0, p , and q match the corresponding functions within Equation (2). We separate out terms dependent and independent of the data x as:

$$\begin{aligned} \Delta(\theta, \theta') &= \sum_{i=1}^N \log \frac{p(x_i | \theta')}{p(x_i | \theta)} - \psi(1, \theta, \theta') \\ &= \Lambda(\theta, \theta') - \psi(1, \theta, \theta'). \end{aligned} \quad (7)$$

A minibatch estimator of Δ , denoted as Δ^* , is

$$\begin{aligned} \Delta^*(\theta, \theta') &= \frac{N}{b} \sum_{i=1}^b \log \frac{p(x_i | \theta')}{p(x_i | \theta)} - \psi(1, \theta, \theta') \\ &= \Lambda^*(\theta, \theta') - \psi(1, \theta, \theta') \end{aligned} \quad (8)$$

Note that Δ and Δ^* are evaluated on the full dataset and a minibatch of size b respectively. The term N/b means $\Delta^*(\theta, \theta')$ is an unbiased estimator of $\Delta(\theta, \theta')$.

²The analysis assumes sampling with replacement although implementations on typical large datasets will approximate this by sampling without replacement.

The key to our test is a smooth acceptance function. We consider functions other than the classical Metropolis test that satisfy the detailed balance condition needed for accurate posterior estimation. A class of suitable functions is specified as follows:

Lemma 2. *If $g(s)$ is any function such that $g(s) = \exp(s)g(-s)$, then the acceptance function $\alpha(\theta, \theta') \triangleq g(\Delta(\theta, \theta'))$ satisfies detailed balance.*

This result is used in Barker [1965] to define the Barker acceptance test.

3.2 Barker (Logistic) Acceptance Function

For our new MH test we use the Barker logistic [Barker, 1965] function: $g(s) = (1 + \exp(-s))^{-1}$. Straightforward arithmetic shows that it satisfies the condition in Lemma 2. It is slightly less efficient than the Metropolis test, since its acceptance rate for vanishing likelihood difference is 0.5. However we will see that its overall sample efficiency is much higher than the earlier methods.

Assume we begin with the current sample θ and a candidate sample θ' , and that $V \sim \mathcal{U}(0, 1)$ is a uniform random variable. We accept θ' if $g(\Delta(\theta, \theta')) > V$, and reject otherwise. Since $g(s)$ is monotonically increasing, its inverse $g^{-1}(s)$ is well-defined and unique. So an equivalent test is to accept θ' iff

$$\Delta(\theta, \theta') > X = g^{-1}(V) \quad (9)$$

where X is a random variable with the logistic distribution (its CDF is the logistic function). To see this notice that $\frac{dV}{dX} = g'$, that g' is the density corresponding to a logistic CDF, and finally that $\frac{dV}{dX}$ is the density of X . The density of X is symmetric, so we can equivalently test whether

$$\Delta(\theta, \theta') + X > 0 \quad (10)$$

for a logistic random variable X .

3.3 A Minibatch Acceptance Test

We now describe acceptance testing using the minibatch estimator $\Delta^*(\theta, \theta')$. From Equation (8), $\Delta^*(\theta, \theta')$ can be represented as a constant term plus the mean of b IID terms $\Lambda_i(\theta, \theta')$ of the form $N \log \frac{p(x_i|\theta')}{p(x_i|\theta)}$. As b increases, $\Delta^*(\theta, \theta')$ therefore has a distribution which approaches a normal distribution by the Central Limit Theorem. We now describe this using an asymptotic argument and defer specific bounds between the CDFs of $\Delta^*(\theta, \theta')$ and a Gaussian to Section 5.

In the limit, since Δ^* is normally distributed about its mean Δ , we can write

$$\Delta^* = \Delta + X_{\text{norm}}, \quad X_{\text{norm}} \sim \tilde{\mathcal{N}}(0, \sigma^2(\Delta^*)), \quad (11)$$

where $\tilde{\mathcal{N}}(0, \sigma^2(\Delta^*))$ denotes a distribution which is approximately normal with variance $\sigma^2(\Delta^*)$. But to perform the test in Equation (10) we want $\Delta + X$ for a logistic random variable X (call it X_{log} from now on). In Bardenet et al. [2015] it was proposed to use Δ^* in a Barker test, and tolerate the fixed error between the logistic and normal distributions.

Our approach is to instead decompose X_{log} as

$$X_{\text{log}} = X_{\text{norm}} + X_{\text{corr}}, \quad (12)$$

where we assume $X_{\text{norm}} \sim \mathcal{N}(0, \sigma^2)$ and that X_{corr} is a zero-mean “correction” variable with density $C_\sigma(X)$. The two variables are added (i.e., their distributions convolve) to form X_{log} . This decomposition requires an appropriate C_σ , which we derive in Section 4. Using X_{corr} samples from $C_\sigma(X)$, the acceptance test is now

$$\Delta + X_{\text{log}} = (\Delta + X_{\text{norm}}) + X_{\text{corr}} = \Delta^* + X_{\text{corr}} > 0. \quad (13)$$

Therefore, assuming the variance of Δ^* is small enough, if we have an estimate of Δ^* from the current data minibatch, we test acceptance by adding a random variable X_{corr} and then accept θ' if the result is positive (and reject otherwise).

If $\tilde{\mathcal{N}}(0, \sigma^2(\Delta^*))$ is exactly $\mathcal{N}(0, \sigma^2(\Delta^*))$, the above test is exact, and as we show in Section 5, if there is a maximum error ϵ between the CDF of $\tilde{\mathcal{N}}(0, \sigma^2(\Delta^*))$ and the CDF of $\mathcal{N}(0, \sigma^2(\Delta^*))$, then our test has an error of at most ϵ relative to the full batch version.

4 CORRECTION DISTRIBUTION

Our test in Equation (13) requires knowing the distribution of X_{corr} . In Section 5, we show that the test accuracy depends on the absolute error between the CDFs of $X_{\text{norm}} + X_{\text{corr}}$ and X_{log} . Consequently, we need to minimize this in our construction of X_{corr} . More formally, let $\Phi_{s_X} = \Phi(X/s_X)$ where Φ is the standard normal CDF³, $S(X)$ be the logistic function, and $C_\sigma(X)$ be the *density* of the correction X_{corr} distribution. Our goal is to solve:

$$C_\sigma^* = \arg \min_{C_\sigma} |\Phi_\sigma * C_\sigma - S| \quad (14)$$

where $*$ denotes convolution. To compute C_σ , we assume the input Y and another variable X lie in the intervals $[-V, V]$ and $[-2V, 2V]$, respectively. We discretize the convolution by discretizing X and Y into $4N + 1$ and $2N + 1$ values respectively. If $i \in \{-2N, \dots, 2N\} = \mathcal{I}$ and $j \in \{-N, \dots, N\} = \mathcal{J}$, then we can write $X_i = i(V/N)$ and $Y_j = j(V/N)$, and

³Hence, Φ_{s_X} is the CDF of a zero-mean Gaussian with standard deviation s_X .

Input : Number of samples T , minibatch size m , error bound δ , pre-computed correction $C_1(X)$ distribution, initial sample θ_1 .

Output: A chain of T samples $\{\theta_1, \dots, \theta_T\}$ from $p(\theta)$;

for $t = \{1, \dots, T\}$ **do**

- Propose a candidate θ' from proposal $q(\theta' | \theta_t)$;
- Draw a minibatch of m points x_i , compute $\Delta^*(\theta_t, \theta')$ and sample variance $s_{\Delta^*}^2$;
- Estimate moments $E|\Lambda_i - \Lambda|$ and $E|\Lambda_i - \Lambda|^3$ from the sample, and error ϵ from Corollary 1;
- while** $s_{\Delta^*}^2 \geq 1$ **or** $\epsilon > \delta$ **do**
 - Draw m more samples to augment the minibatch, update Δ^* , $s_{\Delta^*}^2$ and ϵ estimates;
- end**
- Draw $X_{\text{nc}} \sim \mathcal{N}(0, 1 - s_{\Delta^*}^2)$ and X_{corr} from the correction distribution $C_1(X)$;
- if** $\Delta^* + X_{\text{nc}} + X_{\text{corr}} > 0$ **then**
 - Accept the candidate, $\theta_{t+1} = \theta'$;
- else**
 - Reject and re-use the old sample, $\theta_{t+1} = \theta_t$;
- end**

end

Algorithm 1: Our acceptance test for MCMC.

Table 1: Error (L_∞) in $X_{\text{norm}} + X_{\text{corr}}$ versus X_{log}

N	σ	λ	L_∞ error
4000	0.9	1	1.0e-4
4000	0.8	0.03	5.0e-6

the objective can be written as:

$$C_\sigma^* = \arg \min_{C_\sigma} \max_{i \in \mathcal{I}} \left| \sum_{j \in \mathcal{J}} \Phi_\sigma(X_i - Y_j) C_\sigma(Y_j) - S(X_i) \right|.$$

Now define matrix M and vectors u and v such that $M_{ij} = \Phi_\sigma(X_i - Y_j)$, $u_j = C_\sigma(Y_j)$, and $v_i = S(X_i)$, where the indices i and j are appropriately translated to be non-negative for M , u , and v . The problem is now to minimize $\|Mu - v\|_\infty$ with the density non-negative constraint $u > 0$. We approximate this with least squares:

$$u^* = \arg \min_u \|Mu - v\|_2^2 + \lambda \|u\|_2^2, \quad (15)$$

with regularization λ . The solution is well-known from the normal equations ($u^* = (M^T M + \lambda I)^{-1} M^T v$) and in practice yields an acceptable L_∞ norm.

With this approach, there is no guarantee that $u^* \geq 0$. However, we have some flexibility in the choice of σ in Equation (14). As we decrease the variance of X_{norm} , the variance of X_{corr} grows by the same amount and is in fact the result of convolution with a Gaussian whose variance is the difference. Thus as σ decreases,

$C_\sigma(X)$ grows and approaches the derivative of a logistic function at $\sigma = 0$. It retains some weak negative values for $\sigma > 0$ but removal of those leads to small error. Table 1 shows that the errors between $X_{\text{norm}} + X_{\text{corr}}$ and X_{log} approach single floating precision (about 10^{-7}), and Algorithm 1 describes our procedure. A few points:

- It uses an adaptive step size so as to use the smallest possible average minibatch size. Unlike previous work, the size distribution is short-tailed.
- An additional normal variable X_{nc} is added to Δ^* to produce a variable with unit variance. This is not mathematically necessary, but allows us to use a single correction distribution C_1 with $\sigma = 1$ for X_{corr} , saving on memory footprint.
- The sample variance is denoted as $s_{\Delta^*}^2$ and is proportional to $\|\theta' - \theta\|_2^2$.

5 ANALYSIS

We now derive error bounds for our M-H test and the target distribution it generates. In the most similar prior works, Korattikara et al. [2014] use CLT asymptotic arguments to show that its approximate acceptance test error tends to zero as batch size increases, but no quantitative bounds are given. In Bardenet et al. [2014], explicit bounds are given and depend on

$$C_{\theta, \theta'} = \max_{1 \leq i \leq N} |\log p(x_i | \theta') - \log p(x_i | \theta)|. \quad (16)$$

It is possible to analytically determine $C_{\theta, \theta'}$ for logistic regression, but it is unclear how to do so for more complicated models, and one generally needs to use all $p(x_i | \theta')$ terms⁴ to get $C_{\theta, \theta'}$. In contrast, we use quantitative forms of the CLT which rely on measurable statistics from a *single* minibatch.

In Section 5.1, we present bounds on the absolute and relative error (in terms of the CDFs) of the distribution of Δ^* vs. a Gaussian. We then show in Section 5.2 that these bounds are preserved after the addition of other random variables (e.g., X_{nc} and X_{corr}). It then follows that the acceptance test has the same error bound.

5.1 Bounding the Error of Δ^* from Gaussian

We use the following quantitative central-limit result:

Lemma 3. *Let X_1, \dots, X_n be a set of zero-mean, independent, identically-distributed random variables with sample mean \bar{X} and sample variance s_X^2 where:*

$$\bar{X} = \frac{1}{n} \sum_{i=1}^n X_i, \quad s_X = \frac{1}{n} \left(\sum_{i=1}^n (X_i - \bar{X})^2 \right)^{\frac{1}{2}}. \quad (17)$$

⁴The sample code provided by Bardenet et al. [2014] computes $C_{\theta, \theta'}$ by traversing the entire data, thus providing no performance advantage over the complete test.

This means $t = \bar{X}/s_X$ has an approximate Student's distribution which approaches a Gaussian. Then

$$\sup_x |\Pr(t < x) - \Phi(x)| \leq \frac{6.4E|X|^3 + 2E|X|}{\sqrt{n}}. \quad (18)$$

Proof. See Supplementary Material, Section B. \square

Lemma 3 demonstrates that if we know $E|X|$ and $E|X|^3$, we can bound the error of the normal approximation, which decays as $O(n^{-\frac{1}{2}})$. Making the change of variables $y = xs_X$, Equation (18) becomes

$$\sup_y \left| \Pr(\bar{X} < y) - \Phi\left(\frac{y}{s_X}\right) \right| \leq \frac{6.4E|X|^3 + 2E|X|}{\sqrt{n}} \quad (19)$$

showing that the distribution of \bar{X} approaches the normal distribution $\mathcal{N}(0, s_X)$ whose variance is s_X , measured from the sample.

To apply this to our test, let $X_i = \Lambda_i(\theta, \theta') - \Lambda(\theta, \theta')$, so that the X_i are zero-mean, i.i.d. variables. If instead of all n samples, we only extract a subset of b samples corresponding to our minibatch, we can connect \bar{X} with our Δ^* term: $\bar{X} = \Delta^*(\theta, \theta') - \Delta(\theta, \theta')$, so that $s_X = s_{\Delta^*}$. We can now substitute into Equation (19) and displace by the mean, giving:

Corollary 1.

$$\begin{aligned} \sup_y \left| \Pr(\Delta^* < y) - \Phi\left(\frac{y - \Delta}{s_{\Delta^*}}\right) \right| &\leq \frac{6.4E|X|^3 + 2E|X|}{\sqrt{b}} \\ &= \epsilon(\theta, \theta', b). \end{aligned}$$

Corollary 1 shows that the distribution of Δ^* approximates a Normal distribution with mean Δ and variance $s_{\Delta^*}^2$. Furthermore, it bounds the error with *estimable quantities*: both $E|X|$ and $E|X|^3$ can be estimated as means of $|\Lambda_i - \Lambda|$ and $|\Lambda_i - \Lambda|^3$, respectively, on each minibatch. We expect this will often be accurate enough on minibatches with hundreds of points, but otherwise bootstrap CIs can be computed.

5.2 Adding Random Variables

We next relate the CDFs of distributions and show that bounds are preserved after adding random variables.

Lemma 4. Let $P(x)$ and $Q(x)$ be two CDFs satisfying $\sup_x |P(x) - Q(x)| \leq \epsilon$ with x in some real range. Let $R(y)$ be the density of another random variable y . Let P' be the convolution $P * R$ and Q' be the convolution $Q * R$. Then $P'(z)$ (resp. $Q'(z)$) is the CDF of sum $z = x + y$ of independent random variables x with CDF $P(x)$ (resp. $Q(x)$) and y with density $R(y)$. Then

$$\sup_x |P'(x) - Q'(x)| \leq \epsilon. \quad (20)$$

Proof. See Supplementary Material, Section C. \square

From Lemma 4, we have the following Corollary:

Corollary 2. If $\sup_y |\Pr(\Delta^* < y) - \Phi(\frac{y - \Delta}{s_{\Delta^*}})| \leq \epsilon(\theta, \theta', b)$, then

$$\sup_y |\Pr(\Delta^* + X_{\text{nc}} + X_{\text{corr}} < y) - S(y - \Delta)| \leq \epsilon(\theta, \theta', b)$$

where $S(x)$ is the standard logistic function, and X_{nc} and X_{corr} are generated as per Algorithm 1.

Proof. See Supplementary Material, Section D. \square

Corollary 2 shows that the bounds from Section 5.1 are preserved after adding random variables, so our test remains accurate. In fact we can do better ($O(n^{-1})$ instead of $O(n^{-1/2})$) by using a more precise limit distribution under an additional assumption. We review this in the Supplementary Material, Section E.

5.3 Bounds on the Stationary Distribution

Bounds on the error of an M-H test imply bounds on the stationary distribution of the Markov chain under appropriate conditions. Such bounds were derived in both Korattikara et al. [2014] and Bardenet et al. [2014]. We include the result from Korattikara et al. [2014] (Theorem 1) here: Let $d_v(P, Q)$ denote the total variation distance between two distributions P and Q . Let \mathcal{T}_0 denote the transition kernel of the exact Markov chain, \mathcal{S}_0 denote the exact posterior distribution, and \mathcal{S}_ϵ denote the stationary distribution of the approximate transition kernel.

Lemma 5. If \mathcal{T}_0 satisfies the contraction condition $d_v(P\mathcal{T}_0, \mathcal{S}_0) < \eta d_v(P, \mathcal{S}_0)$ for some constant $\eta \in [0, 1)$ and all probability distributions P , then

$$d_v(\mathcal{S}_0, \mathcal{S}_\epsilon) \leq \frac{\epsilon}{1 - \eta} \quad (21)$$

where ϵ is the bound on the error in the acceptance test.

6 EXPERIMENTS

6.1 Mixture of Gaussians

This model is adapted from Welling and Teh [2011] by increasing the number of samples to 1 million. The parameters are $\theta = \langle \theta_1, \theta_2 \rangle$, and the generation process is

$$\begin{aligned} \theta &\sim \mathcal{N}(0, \text{diag}(\sigma_1^2, \sigma_2^2)) \\ x_i &\sim 0.5 \cdot \mathcal{N}(\theta_1, \sigma_x^2) + 0.5 \cdot \mathcal{N}(\theta_1 + \theta_2, \sigma_x^2). \end{aligned} \quad (22)$$

We set $\sigma_1^2 = 10$, $\sigma_2^2 = 1$ and $\sigma_x^2 = 2$. We fix $\theta = \langle 0, 1 \rangle$. The original paper sampled 100 data points and estimated the posterior. We are interested in performance




Figure 1: The log posterior contours and scatter plots of sampled θ values using different methods.




Figure 2: Minibatch sizes used in Section 6.1's experiment. The axes have the same (log-log scale) range.

on larger problems and so sampled 1,000,000 points to form the posterior of $p(\theta) \prod_{i=1}^{1,000,000} p(x_i | \theta)$ with the same prior from Equation (22). This produces a much sharper posterior with two very narrow peaks. Our goal is to reproduce the original posterior, so we adjust the temperature to $T = 10,000$. Taking logs, we get the target as shown in the far left of Figure 1.

We ran MCMC with our test against the test from Bardenet et al. [2014] and a simplified version of Korattikara et al. [2014]. Both methods use repeated testing which must be discounted to bound overall test accuracy. Bardenet et al. [2014] conservatively partitions the test error, while Korattikara et al. [2014] used a complex dynamic programming heuristic to estimate the overall test error. We did not implement this heuristic and instead set the individual test error equal to the overall test error. Thus our results for Korattikara et al. [2014] should be treated as lower bounds. All methods were initialized with minibatch size 100. For Korattikara et al. [2014] we increment sizes by 100 and set the tolerance $\epsilon = 0.005$. For Bardenet et al. [2014], we increase sizes geometrically with $\gamma = 1.5$ and use parameters $p = 2, \delta = 0.01$. All methods collect 5000 samples using a random walk proposer with covariance $\text{diag}(0.15, 0.15)$, which means the M-H test is responsible for shaping the sample distribution.

Figure 1 shows scatter plots of the resulting θ samples for the three methods, with darker regions indicating a greater density of points. There are no obvious differences, so we measure the similarity between each set of samples and the actual posterior.

We discretize the posterior coordinates into bins with respect to the two components of θ . The probability P_i

of a sample falling into bin i is the integral of the true posterior over the bin's area. A single sample should therefore be multinomial with distribution P , and a set of n (ideally independent) samples is Multinomial(P, n). This distribution is simple and we can use it to measure the quality of the samples rather than use general purpose tests like KL-divergence or likelihood-ratio, which are problematic with zero counts.

For large n , the per-bin distributions are approximated by Poissons with parameter $\lambda_i = P_i n$. Given samples $\{\theta_1, \dots, \theta_T\}$, let c_j denote the number of individual samples θ_i that fall in bin j out of N_{bins} total. We have

$$\log p(c_1, \dots, c_{N_{\text{bins}}} | P_1, \dots, P_{N_{\text{bins}}}) = \sum_{j=1}^{N_{\text{bins}}} c_j \log(nP_j) - nP_j - \log(\Gamma(c_j + 1)). \quad (23)$$

Table 2 shows the likelihoods. To facilitate interpretation we perform significance tests using Chi-Squared distribution (also in Table 2). Scores lie between Korattikara et al. [2014] and Bardenet et al. [2014], but the variance of these values is high and ordering changes depending on the range of samples generated.

Figure 2 shows that the new method dominates in terms of speed and efficiency. The histograms of the (final) minibatch sizes used each iteration show that our method consumes significantly less data; the distribution is short-tailed and the mean is 210, more than an order of magnitude better compared to the other two methods (averages are 15562 and 16857). The sizes correspond to the running times of the methods, since they all consume linear time in data consumed.




Figure 3: Logistic regression performance (accuracy/log likelihood) based on cumulative data usage.




Figure 4: Counts of minibatch sizes in the logistic regression experiment (analogous to Figure 2).

Table 2: Gaussian Mixture Model Statistics

Metric	Ours	Korat.'14	Barde.'14
Equation 23	-1430.0	-1578.9	-1232.7
Chi-Squared	3313.9	3647.7	2444.1

6.2 Logistic Regression

We next test logistic regression for the binary classification of 1s versus 7s on a subset of the MNIST8M dataset, a larger version of MNIST [LeCun and Cortes]. We randomly subsampled 450k training and 192k testing points. We impose a uniform prior on θ and again use a random walk proposer, this time with covariance matrix $0.05I$. The temperature is set at $T = 1000$. We run the three tests for 3000 samples and again use 100 as the starting minibatch size. Method Korattikara et al. [2014] used per-test error of $\epsilon = 0.05$. For Bardenet et al. [2014], we use their recommended discounting factor value of $\gamma = 2.0$. We tried to use the symbolic bound for $C_{\theta, \theta'}$ for Logistic regression from Bardenet et al. [2014], but found it to be too high. Instead we use the empirical $C_{\theta, \theta'}$ from the entire dataset and do not add its computation time when considering runtime, so again our results represent a lower bound for this method.

Figure 3 shows the test log likelihood and prediction accuracy as a function of the cumulative training points processed.⁵ To generate the curves, for each of the

sampled vectors θ_t , $t \in \{1, \dots, 3000\}$, we use θ_t as the parameter for logistic regression. Our minibatch MH test is more efficient in terms of accuracy, achieving convergence roughly twice as fast compared to Korattikara et al. [2014] and about 1.5 orders of magnitude faster than Bardenet et al. [2014], though its log likelihood results lag slightly behind the former algorithm during the very early stages of exploration.

Figure 4 shows histograms of minibatch sizes for the three methods on a log-log scale. With an initial size of 100, we achieve an average minibatch size of 393, far smaller than the averages of the other two methods, which have longer-tailed distributions.

7 CONCLUSIONS

We have derived an M-H test for minibatch MCMC which approximates full data tests. We present theoretical results and experimentally show the benefits of our test on Gaussian mixtures and logistic regression. Directions for future work include running more experiments with a particular focus on controlling variances, testing on neural networks, combining our results with Hamiltonian Monte Carlo methods, providing a recipe for how to use our algorithm (following the framework of Ma et al. [2015]), or integrating parallel MCMC [Angelino et al., 2014, Ahn et al., 2014] concepts.

⁵The curves do not span the same length over the x-axis

since the methods consume different amounts of data.

References

- S. Ahn, A. K. Balan, and M. Welling. Bayesian posterior sampling via stochastic gradient fisher scoring. In *Proceedings of the 29th International Conference on Machine Learning (ICML)*, 2012.
- S. Ahn, B. Shahbaba, and M. Welling. Distributed stochastic gradient MCMC. In *Proceedings of the 31st International Conference on Machine Learning (ICML)*, 2014.
- E. Angelino, E. Kohler, A. Waterland, M. Seltzer, and R. P. Adams. Accelerating MCMC via parallel predictive prefetching. In *Proceedings of the Thirtieth Conference on Uncertainty in Artificial Intelligence (UAI)*, 2014.
- R. Bardenet, A. Doucet, and C. Holmes. Towards scaling up markov chain monte carlo: an adaptive subsampling approach. In *Proceedings of the 31st International Conference on Machine Learning (ICML)*, 2014.
- R. Bardenet, A. Doucet, and C. Holmes. On markov chain monte carlo methods for tall data. *arXiv preprint arXiv:1505.02827*, 2015.
- A. A. Barker. Monte-carlo calculations of the radial distribution functions for a proton-electron plasma. *Australian Journal of Physics*, 18:119–133, 1965.
- V. Bentkus, F. Gotze, and W.R.vanZwet. An edgeworth expansion for symmetric statistics. *Annals of Statistics*, 25(2), 1997.
- S. Brooks, A. Gelman, G. Jones, and X.-L. Meng. *Handbook of Markov Chain Monte Carlo*. CRC press, 2011.
- T. Chen, E. Fox, and C. Guestrin. Stochastic gradient Hamiltonian Monte Carlo. In *Proceedings of the 31st International Conference on Machine Learning (ICML)*, 2014.
- A. Choromanska, M. Henaff, M. Mathieu, G. B. Arous, and Y. LeCun. The loss surface of multilayer networks. In *Proceedings of the 18th International Conference on Artificial Intelligence and Statistics, AISTATS*, 2015.
- W. Gilks and D. Spiegelhalter. *Markov chain Monte Carlo in practice*. Chapman & Hall/CRC, 1996.
- W. K. Hastings. Monte carlo sampling methods using markov chains and their applications. *Biometrika*, 57: 97–109, 1970.
- A. Korattikara, Y. Chen, and M. Welling. Austerity in MCMC land: Cutting the metropolis-hastings budget. In *Proceedings of the 31st International Conference on Machine Learning (ICML)*, 2014.
- Y. LeCun and C. Cortes. MNIST handwritten digit database. URL <http://yann.lecun.com/exdb/mnist/>.
- Y. Ma, T. Chen, and E. Fox. A complete recipe for stochastic gradient mcmc. In *Advances in Neural Information Processing Systems 28*, 2015.
- D. Maclaurin and R. P. Adams. Firefly monte carlo: Exact MCMC with subsets of data. In *Proceedings of the Thirtieth Conference on Uncertainty in Artificial Intelligence (UAI)*, 2014.
- N. Metropolis, A. W. Rosenbluth, M. N. Rosenbluth, A. H. Teller, and E. Teller. Equation of state calculations by fast computing machines. *The Journal of Chemical Physics*, 21, 1953.
- R. M. Neal. MCMC using Hamiltonian dynamics. *Handbook of Markov Chain Monte Carlo*, 54:113–162, 2010.
- Y. Novak. On self-normalized sums and students statistic. *Theory of Probability and its Applications*, 49(2): 336–344, 2005.
- G. O. Roberts and J. S. Rosenthal. Optimal scaling for various metropolishastings algorithms. *Statistical Science*, 16(4):351367, 2001.
- M. Welling and Y. W. Teh. Bayesian learning via stochastic gradient langevin dynamics. In *Proceedings of the 28th International Conference on Machine Learning (ICML)*, 2011.